

# TELOS Talks - Fuzzing

Yonghao Zou

2025.05.22

# TELOS Talks - A Great Plan

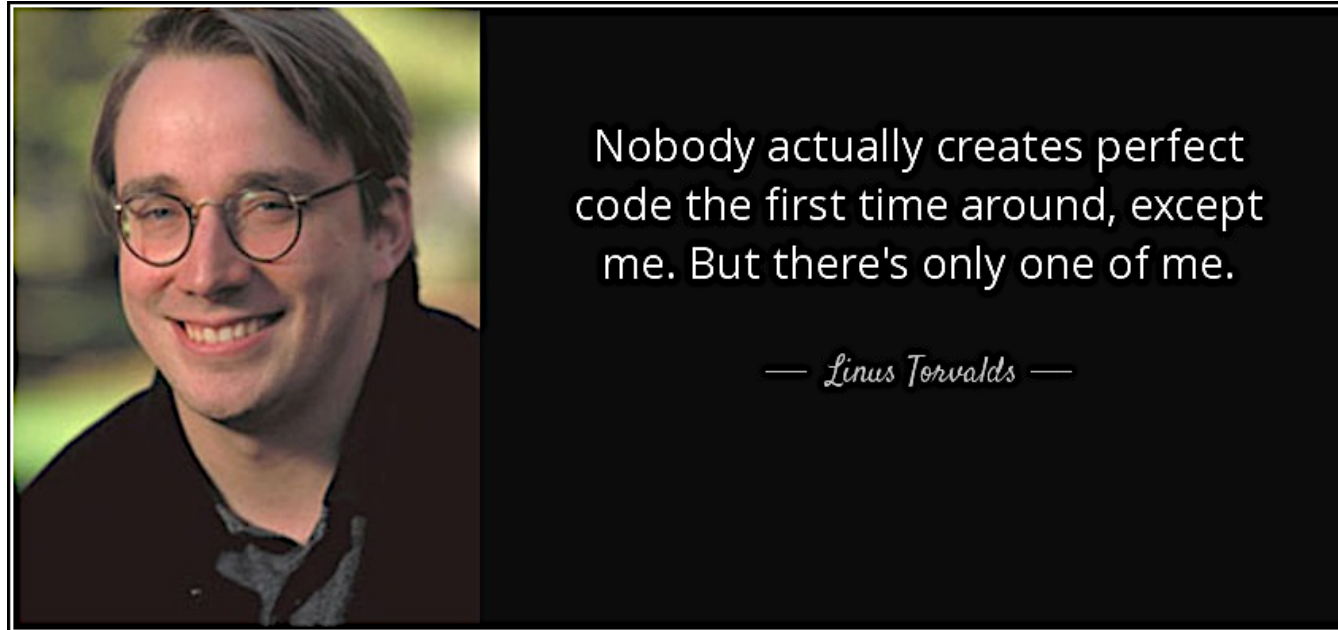
- Why?
  - Learn something new
  - Exchange ideas
  - Add “Talks” to your personal page
- How?
  - Just talk about the areas you know more (than folks in TELOS)
- When?
  - Lunch time

# TELOS Talks - Call for Talks

- Current scheduled talks
  - Introduction to Fuzzing – [Yonghao](#), now
  - Introduction to ML Sys – [Zehua](#), in 2/3 weeks
  - Rust in depth – [Junyang](#), in a month
  - Drawing figures for our paper – [Junyang](#), in 2 months
- Something I can imagine
  - **Intro** to static analysis, LLVM, Verus...
  - **Writing** a symbolic execution engine/verifier, a kernel in Rust, a memory system, a user-level file system, a scheduler using eBPF, **from scratch**
  - File systems, user interrupts, and more system topics **in depth**
  - GPU, FPGA, ML, Tired Memory, ...


# Fuzzing - Background

- Do you write perfect code?



# Fuzzing - Background

- If you are not Linux...
  - Manual
    - Code reviews
    - Unit tests
  - Automatic testing
    - *Fuzzing*
  - Program analysis
    - Static analysis
    - Symbolic execution
  - Formal methods



# Fuzzing - Background

- Pro and cons
  - Unit tests
    - White box
    - Manual effort
  - **Fuzzing**
    - Effective
    - Cannot guarantee the absence of bugs
  - Static analysis
    - Fast, easy to use, static
    - False positive/negative

# Fuzzing - Background

- Pro and cons
  - Symbolic execution
    - Check all possibility of the program
    - Scalability issue
  - Formal methods
    - Prove the absence of certain bugs
    - Manual effort, scalability issue

# Fuzzing - **Blackbox** fuzzing

- You know **nothing or very little** about the program




Program




# Fuzzing - **Blackbox** fuzzing

- Random input
  - E.g., /dev/urandom



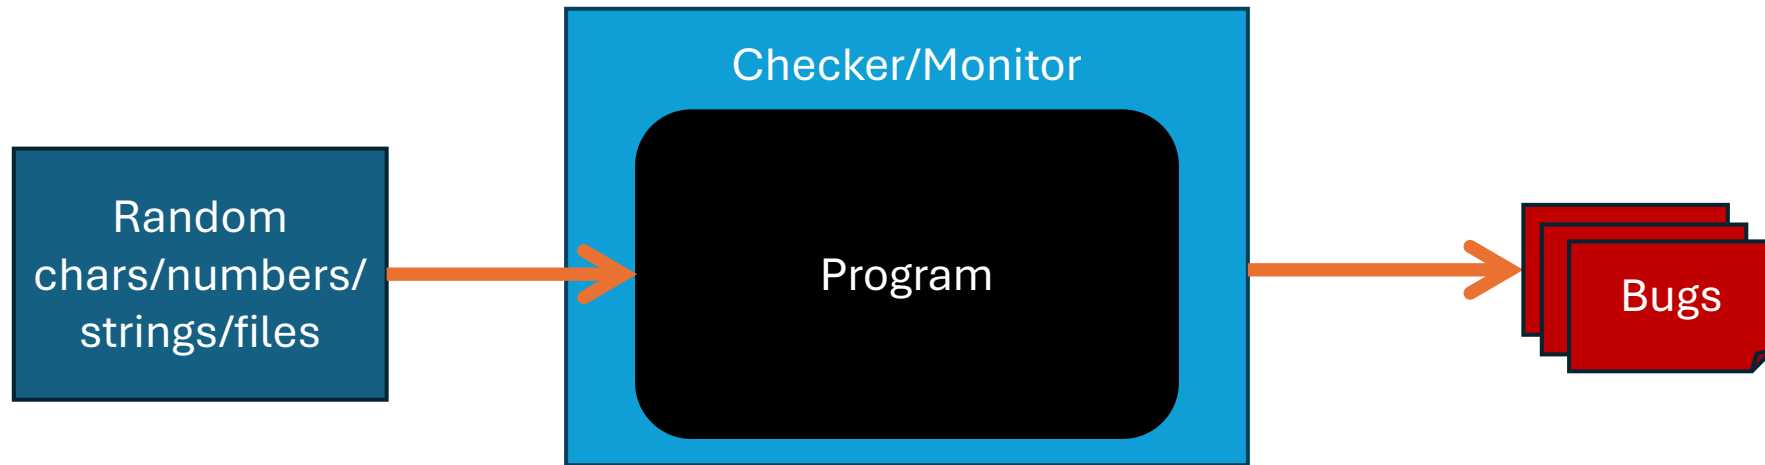
# Fuzzing - **Blackbox** fuzzing

- Checker/Monitor



# Fuzzing - **Blackbox** fuzzing

- Easy, but ineffective



- If we **can know** something about the program, what do we actually want to know?
  - We want to know **whether an input is good or not**

# Fuzzing - Blackbox fuzzing

- **Insight from practice:** If an input is “good”, then a similar input is probably “good”.
  - Input 1: [99, 200]
  - Input 2: [99, 201]

```
if (input[0] < 100) {  
    xxx  
    if (input[1] > 200) {  
        bug_on(...);  
    }  
}
```

# Fuzzing - What is an “good” input?

- A good input can trigger bugs
- **Insight**: more coverage -> more bugs
- Good input: increase coverage

```
if (input[0] < 100) {  
    xxx  
    if (input[1] > 200) {  
        bug_on(...);  
    }  
}
```


# Fuzzing - Coverage

- Line coverage vs. Edge coverage
  - Input 1: [99, 201]
  - Input 2: [99, 200]
  - Is input 2 a good input?

```
if (input[0] < 100) {  
    xxx  
    if (input[1] > 200) {  
        xxx(...);  
    }  
}
```

# Fuzzing - Coverage

- Line coverage vs. Edge coverage
  - Input 1: [99, 201]
  - Input 2: [99, 200]
  - Is input 2 a good input?



# Fuzzing - Mutation

- Mutation-based fuzzing
- E.g, [99, 201]
  - Replace -> [1, 201], [2, 201], [99, 200], [99, 300]
  - Delete -> [99], [201]
  - Duplicate -> [99, 201, 99, 201]
  - Append -> [99, 201, 100, 300]
  - Bitflip
  - ...



# Fuzzing - Generation

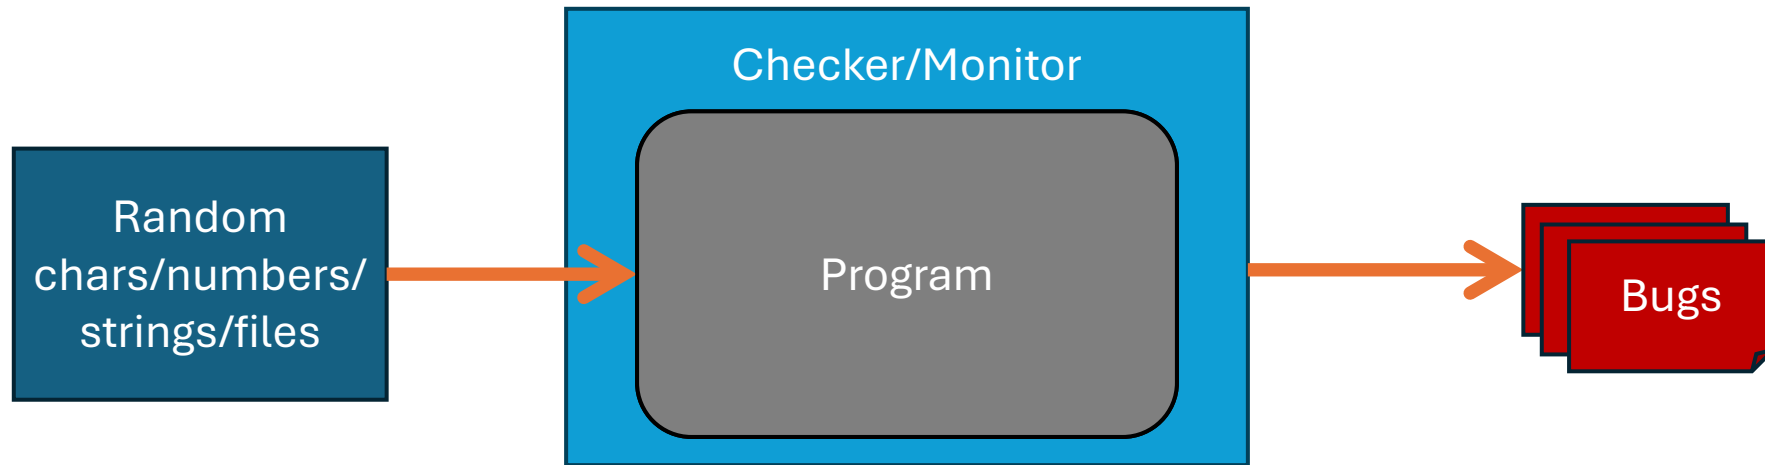
- Generation based fuzzing
- E.g., Network protocol fuzzer usually has a template
  - {IP, Port, Seq number, TCP Options, Timestamp, ...}

```
<!-- A. Local file header -->
<Block name="LocalFileHeader">
  <String name="lfh_Signature" valueType="hex" value="504b0304" token="true" mut
  <Number name="lfh_Ver" size="16" endian="little" signed="false"/>
  ...
  [truncated for space]
  ...
  <Number name="lfh_CompSize" size="32" endian="little" signed="false">
    <Relation type="size" of="lfh_CompData"/>
  </Number>
  <Number name="lfh_DecompSize" size="32" endian="little" signed="false"/>
  <Number name="lfh_FileNameLen" size="16" endian="little" signed="false">
    <Relation type="size" of="lfh_FileName"/>
  </Number>
  <Number name="lfh_ExtraFldLen" size="16" endian="little" signed="false">
    <Relation type="size" of="lfh_FldName"/>
  </Number>
  <String name="lfh_FileName"/>
  <String name="lfh_FldName"/>
  ...
</Block>
```

- In many cases, mutation-based is more effective

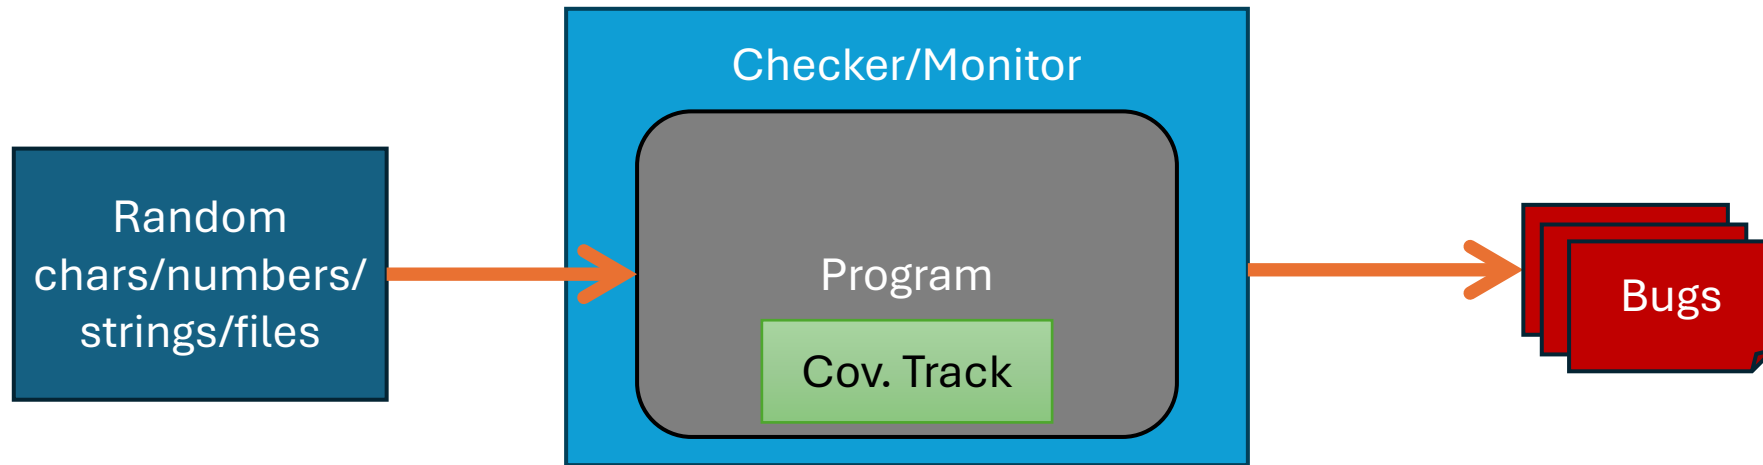
# Fuzzing - Greybox fuzzing

- We do not need to know all, but we can know whether an input is “good”



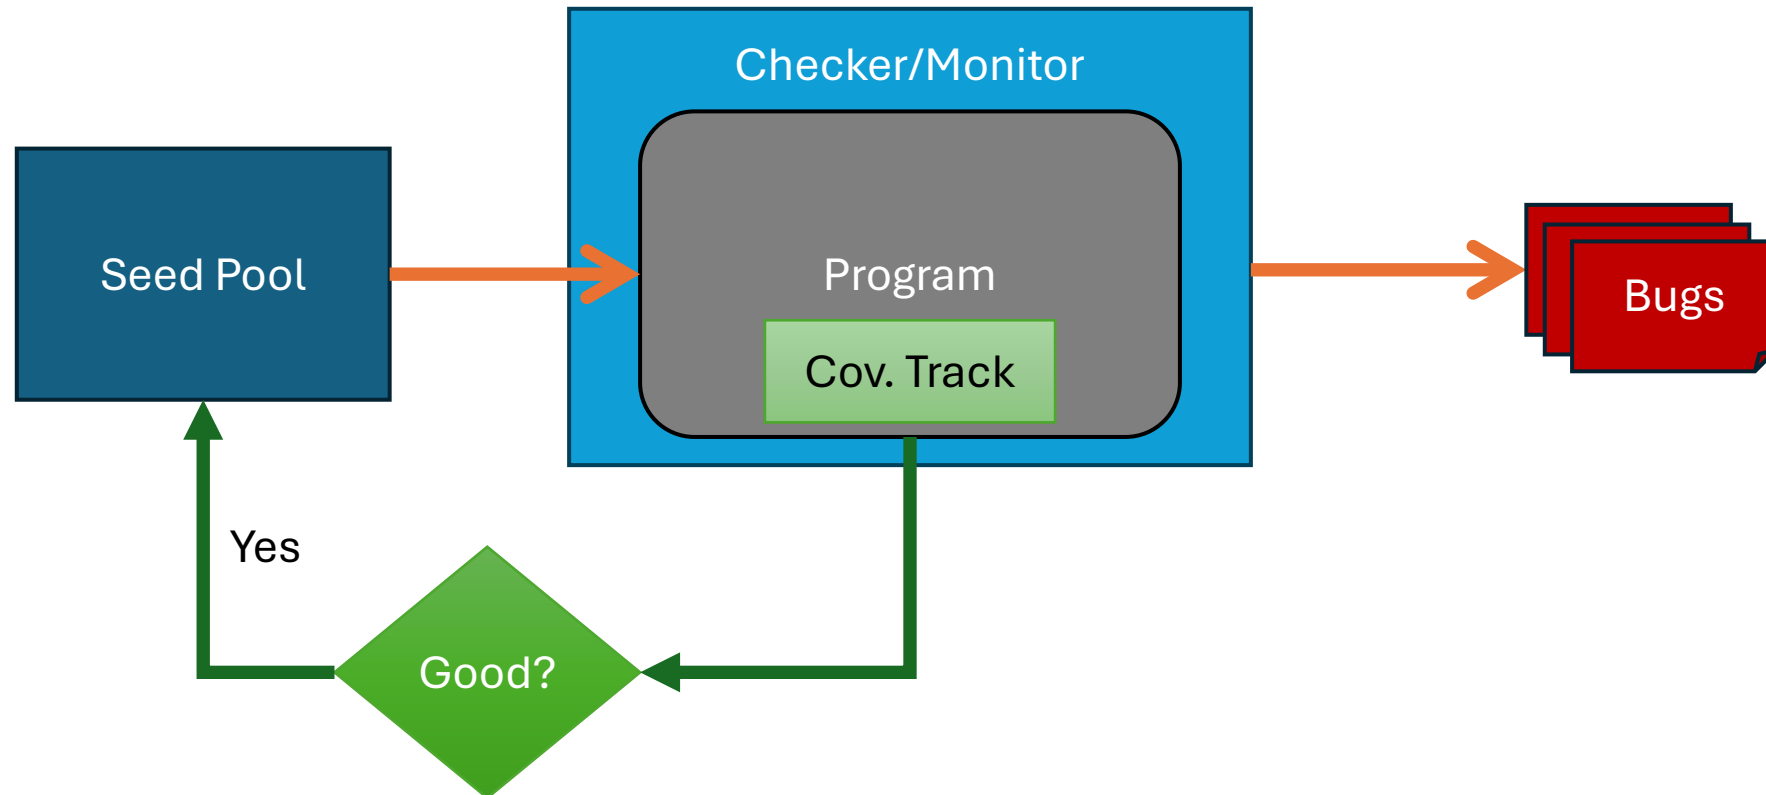
# Fuzzing - Greybox fuzzing

- Coverage tracking



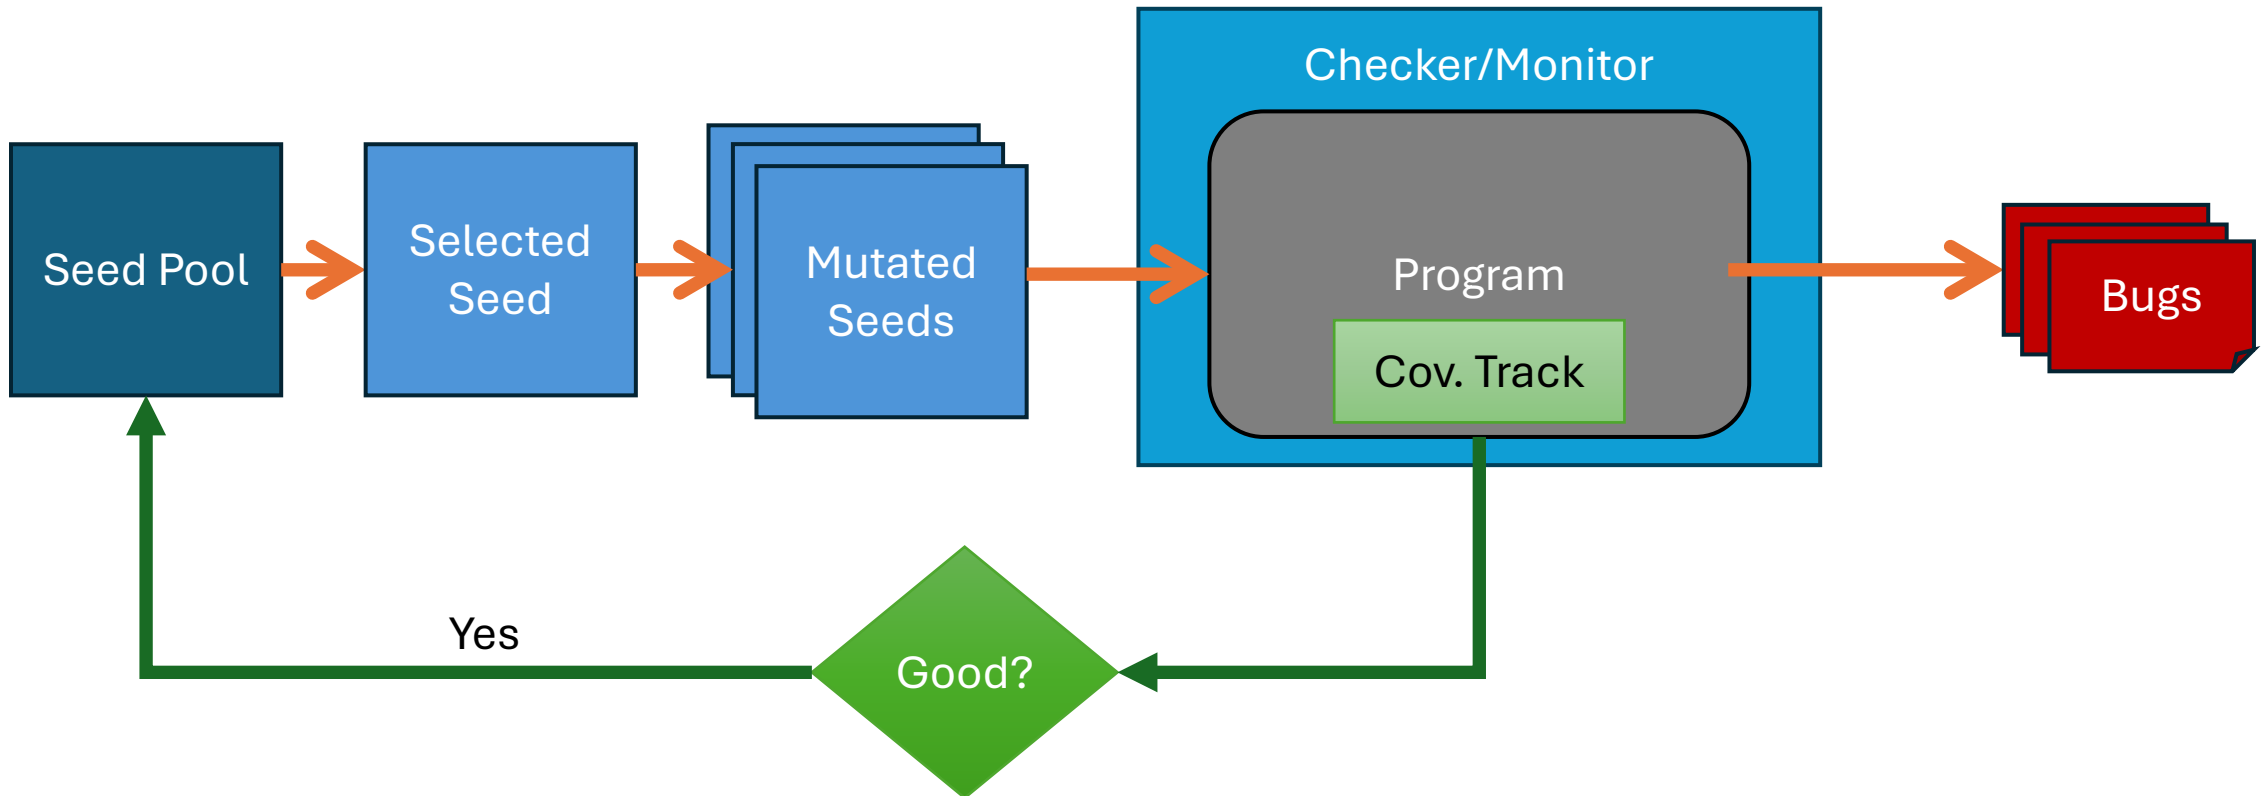
# Fuzzing - Greybox fuzzing

- Input management
  - In fuzzing, we call them “seed”



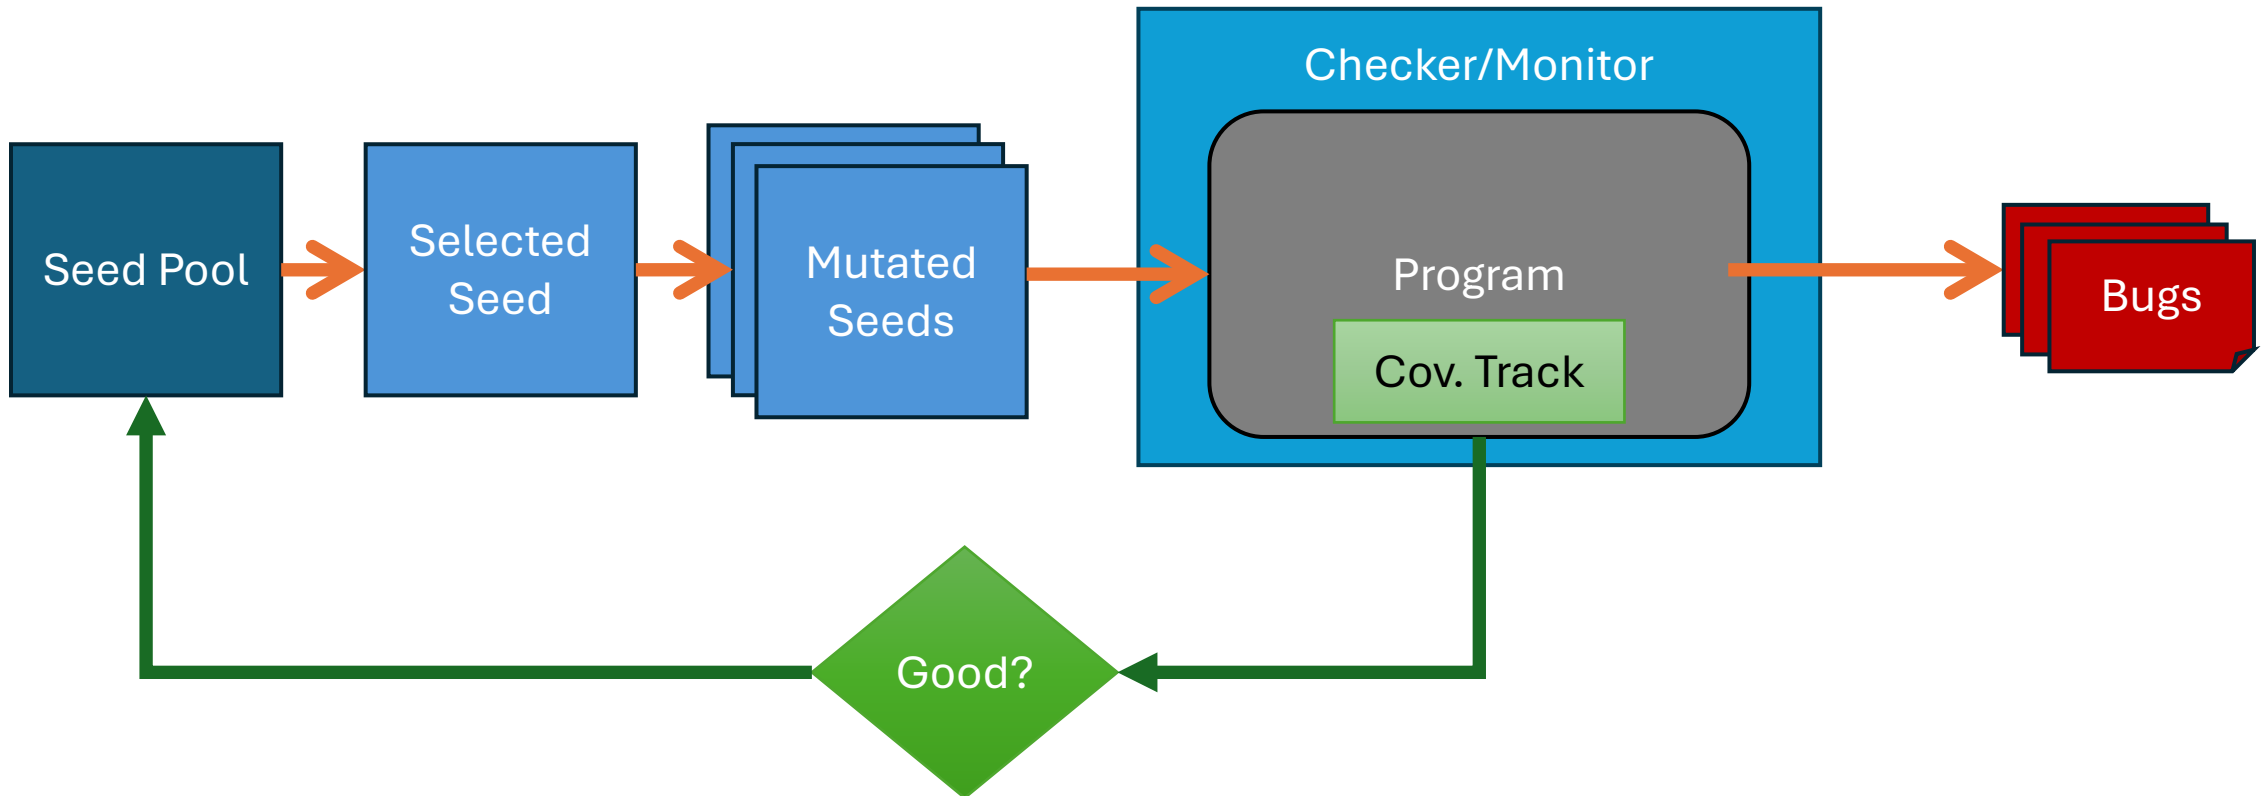
# Fuzzing - Greybox fuzzing

- Seed selection and mutation



# Fuzzing - Greybox fuzzing

- The tool – American Fuzzy Lop (AFL)



# Fuzzing - Greybox fuzzing


- The tool – American Fuzzy Lop (AFL)
  - Filter input with code coverage
  - Mutation-based, few knowledge required
  - Fast: fork-server
  - Checkers
    - ASan
    - TSan
    - ...

# Fuzzing - Greybox fuzzing

- How to get the coverage? Instrumentation

- Insert code for each block
- Must be **short**
- Update the shared memory

```
if (input[0] < 100) {  
    xxx  
    if (input[1] > 200) {  
        xxx(...);  
    }  
}
```




record edge



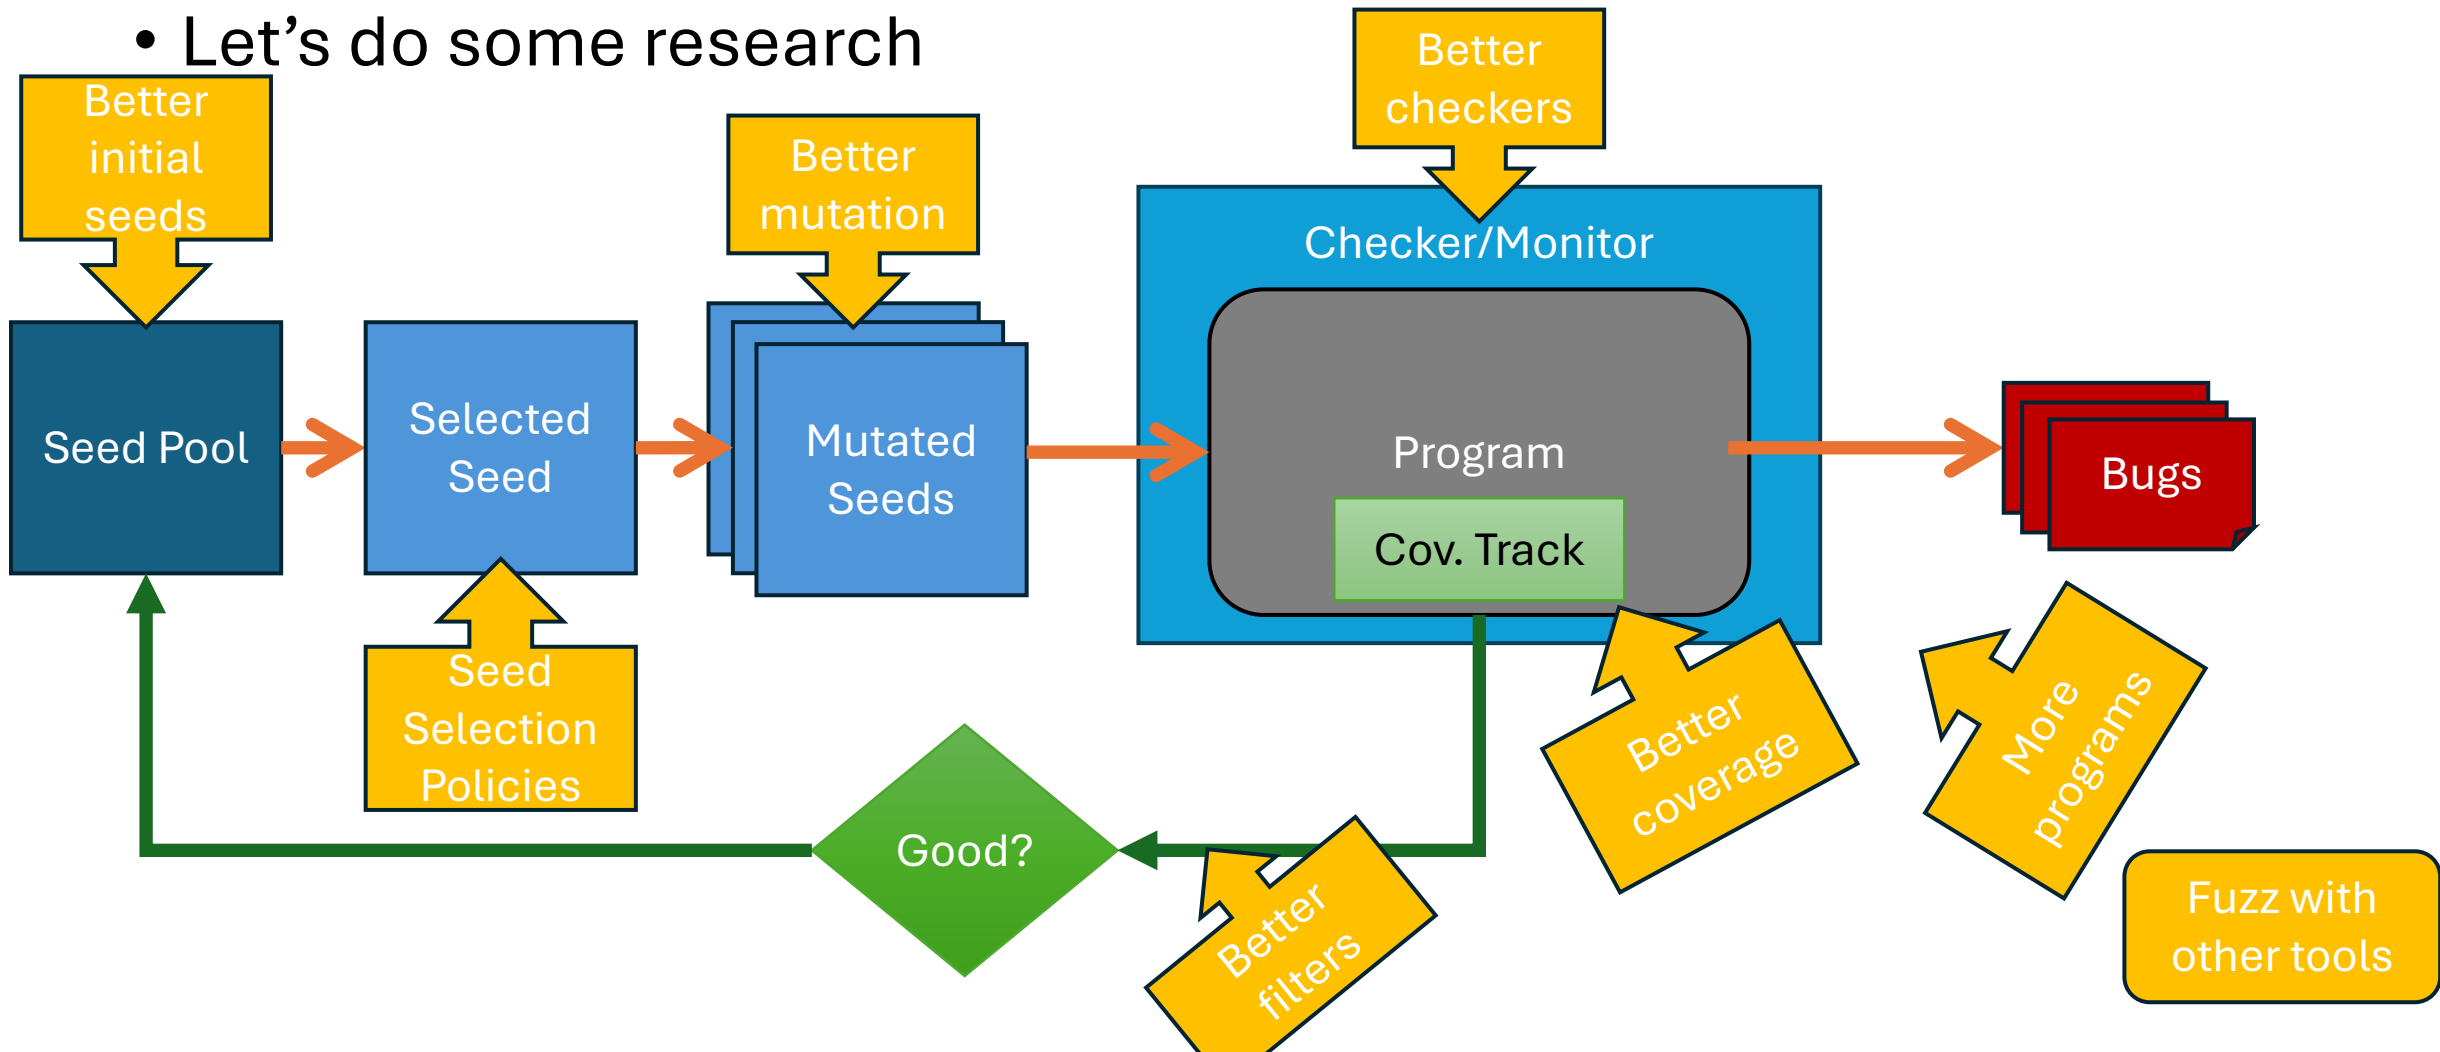
# Fuzzing - Greybox fuzzing

- Fork server
  - To make fuzzing effective, you better make the program run fast ( $>1000$  run/sec)
  - There are some loading/functions before “main”
  - We can pause at main, and fork the process to skip them



# Fuzzing - Research

- Let's do some research



# Fuzzing – Ideas/Cases

- Better initial seeds
  - Use seeds made by experts
  - Generate seeds by symbolic execution -> **concolic** fuzzing (NDSS 16, ISSTA 18, CCS19, CCS 21)
    - What kind of programs is fuzzing ineffective for?

```
if (config->magic != MAGIC_NUMBER)
    puts("Bad magic number");
```

- Any other ideas?
  - 2-Phase fuzz, data analysis, log analysis...

# Fuzzing – Ideas/Cases

- Seed Selection Policies (scheduling)
  - Coverage order
  - Augment order
  - Random
  - Seed minimization (ISSTA 21)
  - CFG analysis (S&P 22)
- Performance varies for different policies (USENIX Sec 21)

# Fuzzing – Ideas/Cases

- Better mutation

```
if (config->magic != MAGIC_NUMBER)
    puts("Bad magic number");
```

- We can know that parts of the numbers match
  - MAGIC\_NUMBER 0xFFFF\_FFFE
  - Current input 0xFFFF\_FFFF
- Mutation based on branch distance (GreyOne, USENIX Sec 20)

# Fuzzing – Ideas/Cases

- Better mutation
  - You can make it complex! (S&P 15)

$\frac{v_c = \text{a concrete value from } src \quad v_a = \{\text{set of input bit positions}\}}{m_c, r_c, m_a, r_a \vdash \text{get\_input}(src) \Downarrow \langle v_c, v_a \rangle}$	INPUT	$\frac{}{m_c, r_c, m_a, r_a \vdash \text{var} \Downarrow \langle r_c[\text{var}], r_a[\text{var}] \rangle}$	VAR
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad v'_c = m_c[v_c] \quad v'_a = m_a[v_c] \hat{\diamond}_b v_a}{m_c, r_c, m_a, r_a \vdash \text{load } e \Downarrow \langle v'_c, v'_a \rangle}$	LOAD	$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad v'_c = \hat{\diamond}_u v_c}{m_c, r_c, m_a, r_a \vdash \hat{\diamond}_u e \Downarrow \langle v'_c, v_a \rangle}$	UNARY-OP
$\frac{m_c, r_c, m_a, r_a \vdash e_1 \Downarrow \langle v_{c1}, v_{a1} \rangle \quad m_c, r_c, m_a, r_a \vdash e_2 \Downarrow \langle v_{c2}, v_{a2} \rangle}{m_c, r_c, m_a, r_a \vdash e_1 \hat{\diamond}_b e_2 \Downarrow \langle v_{c1} \hat{\diamond}_b v_{c2}, v_{a1} \hat{\diamond}_b v_{a2} \rangle}$	BINARY-OP	$\frac{}{m_c, r_c, m_a, r_a \vdash v \Downarrow \langle v, \{\} \rangle}$	CONST
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad \begin{array}{l} r'_c = r_c[\text{var} \leftarrow v_c] \\ r'_a = r_a[\text{var} \leftarrow v_a] \end{array} \quad c' = \text{checkIDS}(c, pc) \quad \boxed{l' = l.\text{add}(\langle v_a, c'.\text{top}() \rangle)} \quad i = s[pc + 1]}{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{var} := e \rightsquigarrow s, m_c, r'_c, m_a, r'_a, \Delta, c', l', pc + 1, i}$	ASSIGN		
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad c' = \text{checkIDS}(c, pc) \quad i = s[v_c]}{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{goto } e \rightsquigarrow s, m_c, r_c, m_a, r_a, \Delta, c', l, v_c, i}$	GOTO		
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle 1, v_a \rangle \quad c = \text{checkIDS}(c, pc) \quad l' = l.\text{add}(\langle v_{a1} \cup v_a, c'.\text{top}() \rangle) \quad i = s[v_{c1}]}{m_c, r_c, m_a, r_a \vdash e_1 \Downarrow \langle v_{c1}, v_{a1} \rangle \quad c' = \text{updateIDS}(c, pc, v_a)} \quad \text{TRUE-COND}$			
$\frac{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{if } e \text{ then goto } e_1 \text{ else goto } e_2 \rightsquigarrow s, m_c, r_c, m_a, r_a, \Delta, c', l', v_{c1}, i}{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle 0, v_a \rangle \quad c = \text{checkIDS}(c, pc) \quad l' = l.\text{add}(\langle v_{a2} \cup v_a, c'.\text{top}() \rangle) \quad i = s[v_{c2}]} \quad \text{FALSE-COND}$			
$\frac{m_c, r_c, m_a, r_a \vdash e_2 \Downarrow \langle v_{c2}, v_{a2} \rangle \quad c' = \text{updateIDS}(c, pc, v_a)}{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{if } e \text{ then goto } e_1 \text{ else goto } e_2 \rightsquigarrow s, m_c, r_c, m_a, r_a, \Delta, c', l', v_{c2}, i}$			
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad c = \text{checkIDS}(c, pc) \quad c' = c.\text{push}(\langle pc + 1, c.\text{pop}() \rangle) \quad i = s[v_c]}{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{call } e \rightsquigarrow s, m_c, r_c, m_a, r_a, \Delta, c', l, v_c, i}$	CALL		
$\frac{m_c, r_c, m_a, r_a \vdash e \Downarrow \langle v_c, v_a \rangle \quad c' = \text{returnIDS}(c, pc) \quad \langle \Delta', l' \rangle = \text{delayedUpdate}(\Delta, l) \quad i = s[v_c]}{s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{ret } e \rightsquigarrow s, m_c, r_c, m_a, r_a, \Delta', c', l', v_c, i}$	RET		
$\frac{m_c, r_c, m_a, r_a \vdash e_1 \Downarrow \langle v_{c1}, v_{a1} \rangle \quad c' = \text{checkIDS}(c, pc) \quad m'_c = m_c[v_{c1} \leftarrow v_{c2}] \quad m'_a = m_a[v_{c1} \leftarrow v_{a1} \hat{\diamond}_b v_{a2}] \quad i = s[pc + 1]}{m_c, r_c, m_a, r_a \vdash e_2 \Downarrow \langle v_{c2}, v_{a2} \rangle} \quad \text{STORE}$			
$s, m_c, r_c, m_a, r_a, \Delta, c, l, pc, \text{store}(e_1, e_2) \rightsquigarrow s, m'_c, r_c, m'_a, r_a, \Delta, c', l, pc + 1, i$			


# Fuzzing – Ideas/Cases

- Better coverage
  - Coverage calculation

$$cur \oplus (prev \gg 1)$$

- Path sensitive coverage (CollAFL, S&P 18)

$$(cur \gg x) \oplus (prev \gg y) + z$$



# Fuzzing – Ideas/Cases

- Better coverage
  - Coverage is often co-designed with seed selection
  - Coverage is often related to the target scenario
- If you want to find memory usage bugs
  - MemLock: Memory Usage Guided Fuzzing
  - The more memory usage, the better



# Fuzzing – Ideas/Cases


- More programs/bugs – **Hot!**
  - You already know OZZ: kernel concurrent bugs
- There's more:
  - Kernel
  - Network protocol (**TCP-Fuzz**)
  - Database
  - Mobile systems (Android, iOS)
  - Embedded systems
  - File systems
  - Distributed systems (**DistFuzz**)
  - Robotic systems (**ROZZ**)

# Fuzzing – Ideas/Cases

- TCP-Fuzz: Detecting Memory and Semantic Bugs in TCP Stacks with Fuzzing (ATC 21)
  - Two-dimensional input: packets and system calls
  - Coverage: state transition coverage

$$BrTran_n = \langle BrCov_n, BrCov_n - BrCov_{n-1} \rangle$$

- Checker: Differential checker




# Fuzzing – Ideas/Cases

- Database Fuzzing
  - You can find how easy the solution is and how vulnerable the system is
  - **Insight:** The result should be identical if two statements are semantically equal
    - `select * from db`
    - `select * from db where 1=1`
    - `select * from db where 1=1 or 1=0`
    - ...
  - EET (OSDI 24) finds 66 bugs in 5 DBMS
  - Yinyang (PLDI 20) finds 1500+ bugs in Z3 and CVC4
  - Try to use similar insights!

# Fuzzing – Ideas/Cases

- Fun papers
  - IJON (S&P 20): manual annotations
  - Play Super Mario better than AFL



# Fuzzing – Ideas/Cases

- Useful tools
  - AFL++: a better AFL that integrates recent research ideas
    - The process of fuzzing libpng



# Fuzzing – Ideas/Cases

- Much more
  - How to fuzz more programs, or things?
    - Compilers
    - Firmware, Drivers
    - Hardware: CPU (S&P 21)
    - Machine learning systems
  - How to make the target program run faster?
    - Checkpoint/restore (Nyx-net, EuroSys 22; DistFuzz, NDSS 25)
  - More types of bugs
  - Integrate with other techniques
    - Static analysis, symbolic execution, formal methods
  - Checkers

# Fuzzing

- Thanks for listening!
- Q&A

# Don't forget

- Add your profile to the lab website!